

Linear Regression

In a Linear Regression problem, we assume that a set of data points can be modeled accurately by a linear function. We further assume that the observed data points contain some unavoidable, random, noise. In these notes we derive the formula for a best fit line under these assumptions.

Given n points $(x_1, y_1) \dots (x_n, y_n)$ and an assumed relation $y = f(x) + \epsilon, \epsilon \sim N(\mu, \sigma)$ we want to find a model $f(x_i) = ax_i + b$ such that the root mean squared error

$$\text{RMSE}(a, b) = \sqrt{\frac{1}{n} \sum_{i=1}^n (\tilde{y}_i - y_i)^2}$$

is minimized. Here we define $\tilde{y}_i = f(x_i)$ and $(\tilde{y}_i - y_i)$ is the *residual* or the error between the observed y_i value and the estimate $\tilde{y}_i$.

First note that if we square the RMSE, the location of the minimum value does not change. Similarly, we can ignore the $\frac{1}{n}$ factor because it does not affect the minimum either. So we can now worry about optimizing the **Residual Sum of Squares** instead.

$$\text{RSS}(a, b) = \sum_{i=1}^n (\tilde{y}_i - y_i)^2$$

is minimized.

RSS is a function of the line parameters a and b only. To minimize it we need to take two partial derivatives and then set both partial derivatives to zero. (This could technically find a maximum – but it's reasonably clear this function has no maximum value because the error can always be increased.)

Take partial derivatives

$$\begin{aligned}
\frac{\partial RSS}{\partial a} &= 2 \sum_{i=1}^n (\tilde{y}_i - y_i) \frac{\partial}{\partial a} (\tilde{y}_i - y_i) \\
&= 2 \sum_{i=1}^n (\tilde{y}_i - y_i) \frac{\partial}{\partial a} (ax_i + b - y_i) \\
&= 2 \sum_{i=1}^n (\tilde{y}_i - y_i) (x_i)
\end{aligned}$$

$$\begin{aligned}
\frac{\partial RSS}{\partial b} &= 2 \sum_{i=1}^n (\tilde{y}_i - y_i) \frac{\partial}{\partial b} (\tilde{y}_i - y_i) \\
&= 2 \sum_{i=1}^n (\tilde{y}_i - y_i) \frac{\partial}{\partial b} (ax_i + b - y_i) \\
&= 2 \sum_{i=1}^n (\tilde{y}_i - y_i) (1)
\end{aligned}$$

And now we set both of the partials equal to zero.

$$\begin{cases} \frac{\partial RSS}{\partial a} = 0 \\ \frac{\partial RSS}{\partial b} = 0 \end{cases} \Rightarrow \begin{cases} \sum_{i=1}^n (\tilde{y}_i - y_i) x_i = 0 \\ \sum_{i=1}^n (\tilde{y}_i - y_i) = 0 \end{cases}$$

Since $\tilde{y}_i = ax_i + b$

$$\sum_{i=1}^n (ax_i + b - y_i) x_i = 0 \Rightarrow a \sum_{i=1}^n x_i^2 + b \sum_{i=1}^n x_i = \sum_{i=1}^n x_i y_i$$

and

$$\sum_{i=1}^n (ax_i + b - y_i) = 0 \Rightarrow a \sum_{i=1}^n x_i + b \sum_{i=1}^n 1 = \sum_{i=1}^n y_i$$

First note that $\sum_{i=1}^n 1 = n$. This is simply a linear system of equations in two unknowns and the coefficients are these messy-looking sums. But they're just numbers – coefficients. We can solve this like any linear system.

The most straightforward way to solve any 2x2 system is using Cramer's rule.

$$a = \frac{\begin{vmatrix} \sum_{i=1}^n x_i y_i & \sum_{i=1}^n x_i \\ \sum_{i=1}^n y_i & n \end{vmatrix}}{\begin{vmatrix} \sum_{i=1}^n x_i^2 & \sum_{i=1}^n x_i \\ \sum_{i=1}^n x_i & n \end{vmatrix}}$$

and solving for b

$$b = \frac{\begin{vmatrix} \sum_{i=1}^n x_i^2 & \sum_{i=1}^n x_i y_i \\ \sum_{i=1}^n x_i & \sum_{i=1}^n y_i \end{vmatrix}}{\begin{vmatrix} \sum_{i=1}^n x_i^2 & \sum_{i=1}^n x_i \\ \sum_{i=1}^n x_i & n \end{vmatrix}}$$

This may look less foreboding if we replace all the sigmas with some new variable names. Let

$$S_x = \sum_{i=1}^n x_i, \quad S_y = \sum_{i=1}^n y_i, \quad S_{xx} = \sum_{i=1}^n x_i^2, \quad S_{xy} = \sum_{i=1}^n x_i y_i$$

Taking determinants,

$$a = \frac{nS_{xy} - S_x S_y}{nS_{xx} - S_x^2} \tag{a}$$

$$b = \frac{S_y S_{xx} - S_x S_{xy}}{nS_{xx} - S_x^2}$$

Interpretation as a ratio of variances

Students of statistics may appreciate the following manipulations

Definition of covariance

$$E(xy) - E(x)E(y) = \text{Cov}(x, y)$$

Definition of variance

$$\text{Var}(x) = E[(x - \mu)^2]$$

Lemma

$$\begin{aligned}\text{Var}(x) &= E[(x - \mu)^2] \\ &= E(x^2) - 2\mu E[x] + E[\mu]^2 \\ &= E[x^2] - 2E[x]^2 + \mu^2 \\ &= E[x^2] - E[x]^2\end{aligned}$$

Manipulating the denominator of equation (a) above,

$$\begin{aligned}n \sum x_i^2 - \left(\sum x_i\right)^2 &= n^2 \left(\frac{1}{n} \sum x_i^2 - \left(\frac{\sum x_i}{n}\right)^2 \right) \\ &= n^2 (E[x^2] - E[x]^2) \\ &= n^2 \text{Var}(x)\end{aligned}$$

And the numerator

$$\begin{aligned}n \sum x_i y_i - \sum x_i \sum y_i \\ &= n^2 \left(\frac{1}{n} \sum x_i y_i - \frac{1}{n} \sum x_i \cdot \frac{1}{n} \sum y_i \right) \\ &= n^2 (E[xy] - E[x]E[y]) \\ &= n^2 (E[xy] - \mu_x \mu_y) \\ &= n^2 \text{Cov}(x, y)\end{aligned}$$

so

$$a = \frac{E[xy] - \mu_x \mu_y}{E[x^2] - \mu_x^2} = \frac{\text{Cov}(x, y)}{\text{Var}(x)}$$